


coarse 891 codes / 9.80 bits + fine 1,225 codes / 10.26 bits = 20.058 bits per bar, both panels ($\lambda = 3$ on the micro block, right panel only)

original — one contextual encoder feeds both halves




bar t 's state is spread across every later bar's token

$\text{sign}(s_t)$ recoverable from bar t 's own token

0.5112

chance 0.50

re-specified — the fine half is per-bar



each bar's fine half is a function of that bar alone

$\text{sign}(s_t)$ recoverable from bar t 's own token

0.9047

interface alone: 0.4995-0.5182, 3/3 FAIL
+ $\lambda=3$ class weight: mean 0.9047, min 0.8839
(restated rule; the run's gate is 0.90 on every dim)